

Contention management for Cloud RAN over an optical ring

Dominique Barth¹ et Maël Guiraud^{1,2} et Yann Strozecki¹

¹David Laboratory, UVSQ

²Nokia Bell Labs France

N-GREEN project has for goal to design a low cost optical ring technology with good performances (dataflow, latency,...) without using expensive end to end connections. We study the compatibility of such a technology with the development of the Cloud RAN, a latency critical application which is a major aspect of 5G deployment. We show that deterministically managing Cloud RAN traffic minimizes its latency without impacting the latency of others traffics.

Mots-clefs : Optical ring, Latency, C-RAN, SDN

TODO: ça marque et dans le titre, ça ne devrait pas

Il faut citer l'ANR NGREEN + n'importe quel truc de charoni qui explique leur architecture NGREEN

1 Introduction

Network providers have to design inexpensive networks while the amount of data and online applications significantly increases. Much of these applications have QoS criteria, like a minimal bandwidth or a maximal latency. The N-GREEN project aims to propose a high performing optical ring while ensuring a minimal cost for providers. The current solutions with good QoS [huawei, demander a dominique C], establishes end to end connection (E2E) between the nodes, which is extremely expensive. The N-GREEN optical ring is designed to ensure good performances and the hardware it requires scales linearly with the number of nodes while E2E scales quadratically making it impractical for more than a few nodes.

In this article, we study a Cloud RAN application in the N-GREEN optical ring. Cloud RAN is one of the major area of development for 5G, that consists in centralizing the computation units or **BaseBand Units** (BBU) of the **Remote Radio Heads** (RRH) in one datacenter. The latency of the messages between the BBU and the RRH is critical. TODO: mettre les citations sur la latence dans la 5G, archi ngreen et concurrents E2E à reprendre de nos précédents articles

TODO: Dire un mot sur la forme du trafic cloud RAN, périodique déterministe, aller vers un data center unique, retour, minimisation de la latence ..., + ne pas déranger BE, décrire le pb

In Sec. 2, we model the optical ring and the traffic flow we study. In Sec. 3, we study the latency when using stochastic multiplexing to manage packets insertions on the ring, with or without priority for C-RAN packets. In Sec. 4, we propose deterministic ways to manage C-RAN packets without buffers and thus with a minimal latency compared to stochastic multiplexing. The best method to do so also improves the latency of best effort packets. TODO: Je sais pas si il faut parler de SDN, qui est un keyword a la mode mais peut etre que l'architecture ngreen nous propose, demander a charoni, Yann : je pense qu'il faut parler de SDN, reprendre peut être des choses de notre article précédent sur Cloud RAN

2 Model of C-RAN traffic over an optical ring

2.1 N-GREEN Optical ring

The unidirectional optical ring is represented by an oriented cycle. The vertices of the cycle represents the nodes of the ring, where traffic arrives. The edges (u, v) of the cycle have an integer weight $\omega(u, v)$ which

represents the time to transmit a unit of information from u to v . We denote by RS (ringsize) the length of the cycle, that is the sum of $\omega(u, v)$ for all arcs (u, v) . Each arc (u, v) can be seen as a sequence of $\omega(u, v)$ **containers**, of capacity C , expressed in Bytes. The containers are filled by the nodes. The node u can put some data in the first container of the arc (u, v) if it is free. The dynamic of the network is simple: at each unit of time data contained in a container go to the next. The ring follows a **broadcast and select with emission release policy**. When a container is filled by some node u , it is freed when it comes back at u after a whole round trip.

2.2 C-RAN traffic

We study a C-RAN traffic emitted by each RRH. The data from an antenna i arrives into a specific node u at time m_i , called **offset**, then it is routed to its BBU at node v through the ring. Then after some computation time (which is here supposed to be zero to simplify the model), an answer is sent back from the BBU to the RRH at time $t_i = m_i + \omega(u, v)$, called **BBU-offset**. The same quantity of data is emitted by an antenna at each period P . In our model all BBUs are in the same data-center, thus they are all attached to the same node of the ring. In the following, we will denote by n the number of antennas. The **latency** of a message is the time it waits into a node before being put in a container.

The antenna are linked to the nodes of the ring through an electronic interface of bit rate $RBps$. The ring has a larger rate of $FRBps$. The integer F is called the **acceleration factor** between the electronic and the optical domains. A node aggregates the data received on the electronic interface during F unit of times to create a PDU of size C , that is which can fills one container of the ring. As a consequence, we consider that each period, an antenna emits a PDU of size C each F unit of times during a time ET .




Figure 1: Opto-electronic interface.

TODO: Améliorer le dessin: un seul schéma avec les notations et horizontal pour pas perdre de place. Éventuellement faire une représentation du ring et en zoom faire l'interface

TODO: Décrire ici les degrés de liberté -ζ quand ça part + la latence

2.3 Best effort traffic

The traffic on the ring is also composed of a best effort traffic, corresponding to the internet flow. We model the arrival of best effort messages by a temporal law that gives the distribution of the arrival between two arrivals of containers with best effort message in a node. This distribution is obtained by studying the impact of the fill rate and the maximum time to wait before sending a container on a contention buffer. This contention buffer is filled by some batch arrivals of data, and creates the containers. [ref].

2.4 Parameters of the study

The experiment of this subsections are made following the same parameters, until sec 4.2. We set $K = 10$ channels, some electronic interface of rate $D = 10\text{Gbps}$ and $Z = 12500\text{B}$ so, $Z/KD = 1\mu\text{s}$. The optical ring has a rate $KD = 100\text{Gbps}$, while a slots has a duration of $1\mu\text{s}$. The length of the ring is to 20km in our simulations, this means that the ring has a length of $RS = 100$ slots. We consider some BE SDUs of 125B , and we parametrized BE PDU generator of sec ?? such that each node load the network with an average of

Number of channels K	10
Rate of an electronic interface D	10 Gbps
Container size Z	12,500 B
Slot duration Z/KD	1 μ s
Optical ring rate KD	100 Gbps
Length of the ring RS	100 slots
BE SDU size	125B
C-RAN PDU size	12,500B = Z
Emission time ET	500 slots
Period P	1,000 slots
Number of RRH	5
Number of nodes n	5

Figure 2: Parameters of the following experiments based on N-GREEN

3,5% of BE traffic. We set $ET = 500$ slots, that correspond to a C-RAN traffic of 5Gbps. The period P is to 1ms, that is 1000 slots. We set the number of RRH to $n = 5$. We take 5 nodes and one of them is the datacenter. The antennas are apportioned between the others node. With those parameters, the network is theoretically loaded to 67.5%.

3 Performance evaluation of the N-GREEN optical ring

In a first time, we will look at the performance of the N-GREEN optical ring with an opportunistic insertion policy. When a node wants to send a PDU on the ring, it uses the first available container. We study the performance of the ring through two methods; a full opportunistic method, and a method that prioritize the C-RAN PDUs. Since we only consider some PDUs arrivals, those two methods are some variants of the management of the PDUs in the node. For those two methods, the offsets m_r of the routes are chosen randomly in the period, i.e. we do not choose the first time at which the RRH start to emit.

3.1 Full opportunistic method

We first want to look at a simple configuration in which we do not manage the PDUs (BE or C-RAN) arriving into a node. At their arrival, the BE and C-RAN PDUs are buffered in the contention buffer, and this buffer is emptied following the FIFO rule.

The following experiment shows us the performance this **full opportunistic method** with some parameters issued from N-GREEN. On fig 3.1, one can see the cumulative distribution of the latency of the C-RAN and BE PDUs. Those results are computed on 100 instances of optical rings that have run during 10,000,000 slots. The measures are taken after 3,000 slots of simulation, to let the ring initialize its behavior. The parameters of the simulations are stated in sec 2.4.

As we can see on fig 3.1, the latency for the different kind of PDU is notably similar because we do not manage anything. Furthermore, we can remark that about 10% of the C-RAN messages wait over 100 μ s, that is already critical for C-RAN **TODO: ref.**

3.2 C-RAN priority method

In order to improve the C-RAN PDUs latency, we imagined the following solution, called **C-RAN priority method**. Each nodes have two contention buffers. One for the BE PDUs, and one for the C-RAN PDUs. When a node is able to send a PDU on the ring, i.e. its available container is free, the container is filled with the C-RAN PDUs first.

With this method, one can imagine improve the latency of the C-RAN PDUs. Fig 3.2 shows us the performance of the C-RAN priority algorithm for different kinds of traffics. The parameters of the simulations still the same, stated in sec 2.4.

As expected, the latency of the C-RAN PDUs is decreased compared to the full opportunistic method. Nevertheless, there is still 10% of the C-RAN PDUs that have more than 50 μ s of latency, and the best effort

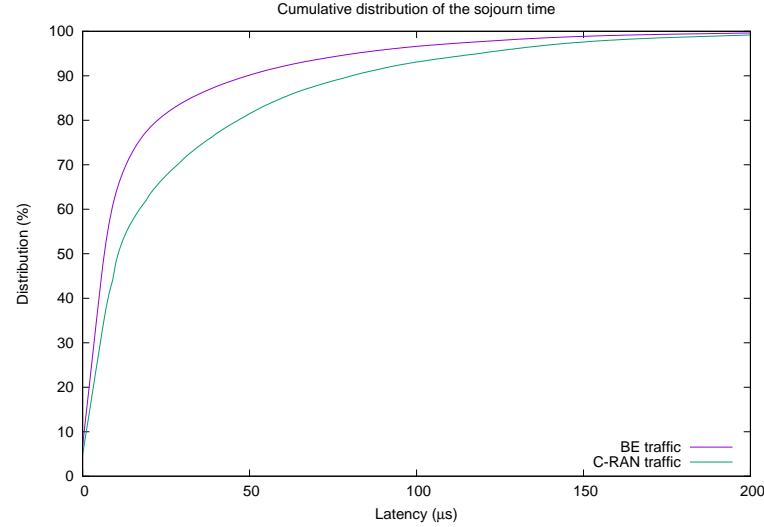


Figure 3: Full opportunistic method performances.

PDU's are penalized, indeed, close to 10% of the BE PDU's awaits more than $150\mu s$ before being sent in the ring, versus close to 0% with the previous method.. Thus, we propose a deterministic approach in which we set the offsets of the RRH and we reserve the containers needed for the C-RAN PDU's. **TODO: il manque une partie sur le travail de karim,**

4 Deterministic approach for 0 latency on C-RAN

In this section We present three algorithms that allows the C-RAN PDU's to have 0 latency in the insertion buffers by establishing a static reservation of the slots in the period. Those algorithm are **Centralized** and requires a global supervisor of the network to be implemented. **TODO: phrase + ref sdn?**

Slot reservation To decrease the C-RAN messages latency, one plan to reserve the container for the nodes. Indeed, if a node i needs to send a PDU at time t , a container needs to be reserved since the previous lap of the container, i.e. at time $t - RS$ by i . In this situation, if the container is already taken by a PDU which has been sent by a node j , j will remove the PDU of the container during the next lap of the ring, and no other nodes are able to write in this container before it comes back to the node i a date t . By choosing the offsets of the route and reserving the corresponding containers, we allow the C-RAN PDU's to have 0 latency in the nodes. Though, it impacts the best effort PDU's, by avoiding them to use free reserved containers. It is equivalent to increase the load of the network, without increasing the number of PDU's carried by the network.

RRHs synchronisations This slot reservation implies that the RRHs are synchronized with the ring. The following proposed algorithms set a sequence of reserved slots for each RRH in each period. A technical application implies to simultaneously ensure the reservation of the slot to the given dates, and the synchronization of the RRH with those reserved slot. One must ensure that the C-RAN PDU arrives in the contention buffer of the node just before the reserved container.

Unlike the previous methods with opportunistic insertion policy, we now want to choose a the offsets for the routes. Since all the forward routes have the same last vertex v (the datacenter), which is also the first vertex of all the backward routes, we choose to manage the collision in this vertex only. The purpose is then




Figure 4: C-RAN priority method.

to determine the time BBU offset of the routes. Once the BBU-offsets are set, for all routes i , it is easy to compute $m_{r_i} = t(v, r_i) - \omega(u, v)$.

4.1 *N-GREEN parameters adapted algorithms*

An RRH can not emit a message more often than once slot every K slots. Remember that **macro slot** is a duration of K slots, during which each RRH can emit at most once. Thus, we consider the simple case in which there is less C-RAN traffic than slots in a macro slot. Indeed, in that case, we only manage the collision in one macro slot, i.e. we assign a C-RAN traffic to each slot of the macro-slot. Since a BBU always instantly emits an answer to its RRH, for each RRH, a group of two following containers (called **bunch**) can be reserved. Thus, we first study the case $K \geq 2 \times n$. As a reminder, n is the number of couple BBU-RRH.

Because there is at most as much couples BBU-RRH as bunches, we want to assign one RRH to one bunch. We propose the following **naïve reservation algorithm**. For each RRH $i, i \in \{0, \dots, n-1\}$, set $t(v, r_i) = 2 \times i$ with v the vertex of the BBUs. The model automatically fixes the offset of the backward route $\rho(r_i)$ to $m_{\rho(r_i)} = 2 \times i + 1$. This is the reason why we do not care about the backward routes anymore, and we only deal with the BBU-offsets of the forward routes..

We kept the same parameters than in the two previous experiments (see 2.4) and we made an experiment to show the impact of this algorithm on the BE traffic. Fig 7 shows us the distribution of latency for BE PDUs with this algorithm.

As we can see, the BE traffic is highly penalized by this algorithm, indeed, while almost 100% of the BE PDUs have a latency to less than $200\mu s$ with the full opportunistic method (see fig 3.1), here, more than half of the BE PDUs are buffered more than $200\mu s$. Those performances are explained by the structure of the algorithm. The messages are grouped together, and with our parameters, $K = 2 \times n$. It creates a long sequence of slots during which there is no free containers in the ring.

To improve the BE PDUs latency, we propose a **split reservation algorithm** that balances the load of the CRAN traffic over the period. Instead of giving some close BBU-offsets for all routes i , with v the datacenter, we now uniformly distribute the different BBU-offsets in the period. Each couple BBU-RRH is still assigned to its bunch, but we split the BBU-offsets of $\frac{P}{n}$. Also, if $2 \times n < K$, the RRH are balanced over the bunches by the same way, we try to give a bunch to a RRH each $\frac{K}{2 \times n}$ bunches.

On fig 9 one can observe the performance of this algorithm on the BE PDUs latency. Once again, to have




Figure 5: C-RAN priority method performances.

a reference point, we take the same parameters for the simulation.

As we can see, the latency of the BE PDUs is highly better with this algorithm. Indeed almost 100% of CRAN PDUs have a latency of less than $50\mu s$ which is better than the full opportunistic method. Thus, we decreased the latency of every kind of traffic.

4.2 With more antennas

In this section, we study the case $n > \frac{K}{2}$, we want to assign more than one RRH to a bunch. Let us consider a particular behavior of the ring. Instinctively, we can think that if we have enough space in the period for many traffics on the same frequency that is, $ET \leq \frac{P}{2}$, we can assign several couples BBU-RRH to the same bunch.

Let us take $ET = \frac{P}{2}$ and $n = \frac{K}{2} + 1$. We assign a bunch to each $\frac{K}{2}$ first RRH and we have only one RRH left. Then, we only look at one bunch, on which we want to add the new couple BBU-RRH.

We take a ring with two nodes A and B . We simplify the model to only look at the behaviour of one bunch, thus, both of the nodes alternatively emits some containers every slots, during ET slots in the period P . We assume that A emits a time 1 during ET slots. This container arrives in B at $\omega(A, B)$, that is $[t(B, r_A)]_{P, K} = [\omega(A, B); \omega(A, B) + ET]$. Then, B can immediately emit its containers at time $\omega(A, B) + ET$. This container from B arrives at A at time $\omega(A, B) + ET + \omega(B, A) = ET + RS$, and thus, $[t(A, r_B)]_{P, K} = [ET + RS; ET + RS + ET] = [ET + RS; RS + P] = [ET + RS; RS] \mod P$. Since $[t(A, r_A)]_{P, K} = [1; ET]$, there is a collision at the interval $[1; RS]$ in the period between the two traffics from A and B .

As we have just noted, if several RRH share the same bunch, the period must have an additional budget in addition to the time needed to send the messages. Let us study the size of the additional budget.

Proposition 1. *If several RRH shares a bunch, it is possible set the offsets of the routes such that there is no collisions between the messages if $P \geq n \times ET + RS$, with n the number of RRHs.*

Proof. We seek to show that proposition by induction.

Base case: For $n = 2$, the proof of the previous example shows us that it is possible to assign two RRH on the same bunch with $P = 2 \times ET + RS$.

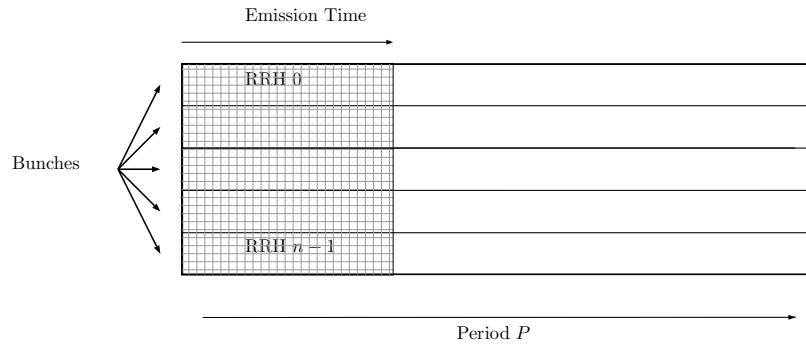


Figure 6: Grouped assignment of the bunches to the RRH.




Figure 7: Impact of the naive reservation algorithm on the Best effort PDUs.




Figure 8: Splited assignment of the bunches to the RRH.




Figure 9: Impact of the splitted reservation algorithm on the Best effort.

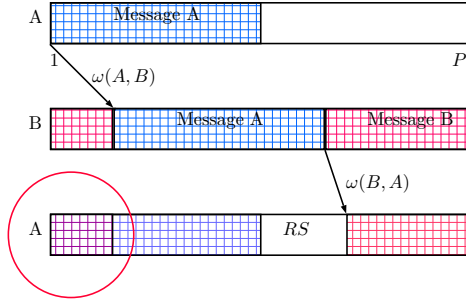


Figure 10: Collision between two routes on the same bunch when the period is too small.

Induction step: We consider that we have a bunch shared by n RRH, with a period of size $P = n \times ET + RS$. We want to show that if we want to add one RRH to the bunch, we need a period of size at least $P = (n+1) \times ET + RS$. If we look at the sequence period of the bunch with n RRH. We can observe that this sequence is the same in every nodes, shifted of the length of the routes between two nodes. For instance, if we take two nodes A and B , the sequence in A is the same that the sequence in B shifted of $\omega(A, B)$.

Thus, if we add ET slots in this bunch, the periodic emissions of all others nodes are not impacted, there is just a new interval in the sequence.

Note that a station can send some traffic from more than one RRH, the figures 11 and `feffig:proofperiod2` are just an example, but this proof stays the same. \square

Considering those results, we propose an algorithm to allow several RRH to share the same bunch, as long as $P \geq n \times ET + RS$. Since we need RS slots on a each bunch with at least an RRH to assign another RRH to this latter, it is interesting to fill as far as possible the bunches. We then compute the number of bunches needed to carry the number of RRH n . Then, the algorithms tries to balance those bunches into the macro-slot, i.e. it avoids to group the used bunches together. The used bunches have RS or more free slots in the period. Thus, the algorithm subsequently balance BBU-offsets of the first RRH of each bunch, with




Figure 11: An example of sequence with $n = 3$.




Figure 12: The same sequence with $n = 4$.

the same idea than the splitted algorithm. We call this algorithm the **full repartition algorithm**.




Figure 13: Organization of the bunches with the balancing algorithm.

We want to observe the impact of this algorithm on BE traffic. Since our previous parameters did not allow several RRH to share a bunch, we now set $ET = 200$, that correspond to a C-RAN flow of 2Gbps. This is not out of context since the exact split (the C-RAN split is the degree of centralization of the computation units in the cloud) of the C-RAN is not fully determined yet [?]. **TODO: la ref** To keep a load similar to the load in the previous experiment, we set the number of RRH to $n = 12$. The others parameters are all the one described in the sec 2.4. The following experiment shows us the impact of the balancing algorithm on the BE traffics. We chose to compare those results to a simulation with the same parameters but with the full opportunistic method of sec 3.1.

As fig 14 shows, the full repartition algorithm ensures a lower latency to BE PDUs than the full opportunistic method, which gave the best result for the BE traffic. Indeed, since we remove the random aspect of the C-RAN PDUs generation, we balanced the load over the period an then increased the BE performances while we totally removed the C-RAN PDUs latency.




Figure 14: Balancing algorithm performance on the N-GREEN optical ring.